

```
# Importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, roc_auc_score, roc_curve

# Load the dataset
data = pd.read_csv('/content/Preprocessed_Missing_dataset.csv')
# Display the first few rows of the dataset
data.head()
```



	ID	Customer_ID	Month	Name	Age	SSN	Occupation	Annual_Income	Monthly_Inhand_Salary	Nu
0	0x1602	CUS_0xd40	January	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	
1	0x1603	CUS_0xd40	February	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	
2	0x1604	CUS_0xd40	March	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	
3	0x1605	CUS_0xd40	April	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	
4	0x1606	CUS_0xd40	May	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	

5 rows × 28 columns



```
# Display summary statistics
print(data.describe())

# Check data types and info
print(data.info())

# Check for unique values in each column
print(data.nunique())
```



```

22 Payment_of_Min_Amount      100000 non-null object
23 Total_EMI_per_month        100000 non-null float64
24 Amount_invested_monthly    100000 non-null float64
25 Payment_Behaviour          100000 non-null object
26 Monthly_Balance            100000 non-null float64
27 Credit_Score                100000 non-null object

```

```
dtypes: float64(9), int64(7), object(12)
```

```
memory usage: 21.4+ MB
```

```
None
```

```
ID                                100000
```

```
Customer_ID                      12500
```

```
Month                             8
```

```
Name                             10139
```

```
Age                               43
```

```
SSN                              12500
```

```
Occupation                       15
```

```
Annual_Income                   12488
```

```
Monthly_Inhand_Salary           13235
```

```
Num_Bank_Accounts                12
```

```
Num_Credit_Card                  1179
```

```
Interest_Rate                    1750
```

```
Num_of_Loan                      10
```

```
Type_of_Loan                     6260
```

```
Delay_from_due_date              68
```

```
Num_of_Delayed_Payment           708
```

```
Changed_Credit_Limit             3716
```

```
Num_Credit_Inquiries             18
```

```
Credit_Mix                       3
```

```
Outstanding_Debt                 12203
```

```
Credit_Utilization_Ratio         100000
```

```
Credit_History_Age               414
```

```
Payment_of_Min_Amount            3
```

```
Total_EMI_per_month              12191
```

```
Amount_invested_monthly          97607
```

```
Payment_Behaviour                 6
```

```
Monthly_Balance                  99759
```

```
Credit_Score                      3
```

```
dtype: int64
```

```
# Plotting histograms for all features
```

```
data.hist(bins=30, figsize=(15, 10))
```

```
plt.tight_layout()
```

```
plt.show()
```


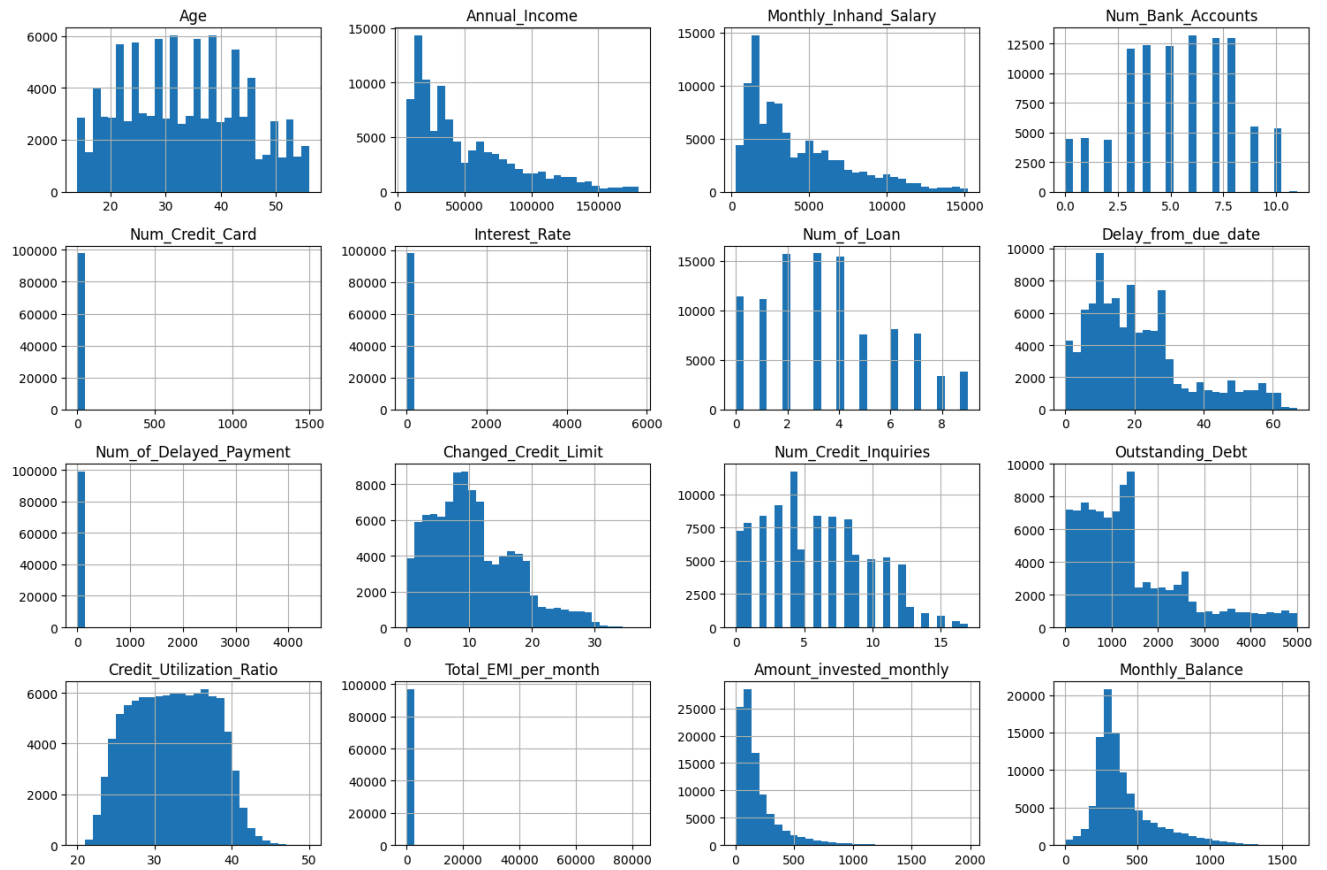
```
# Correlation heatmap
```

```
numeric_data = data.select_dtypes(include=[np.number])
```

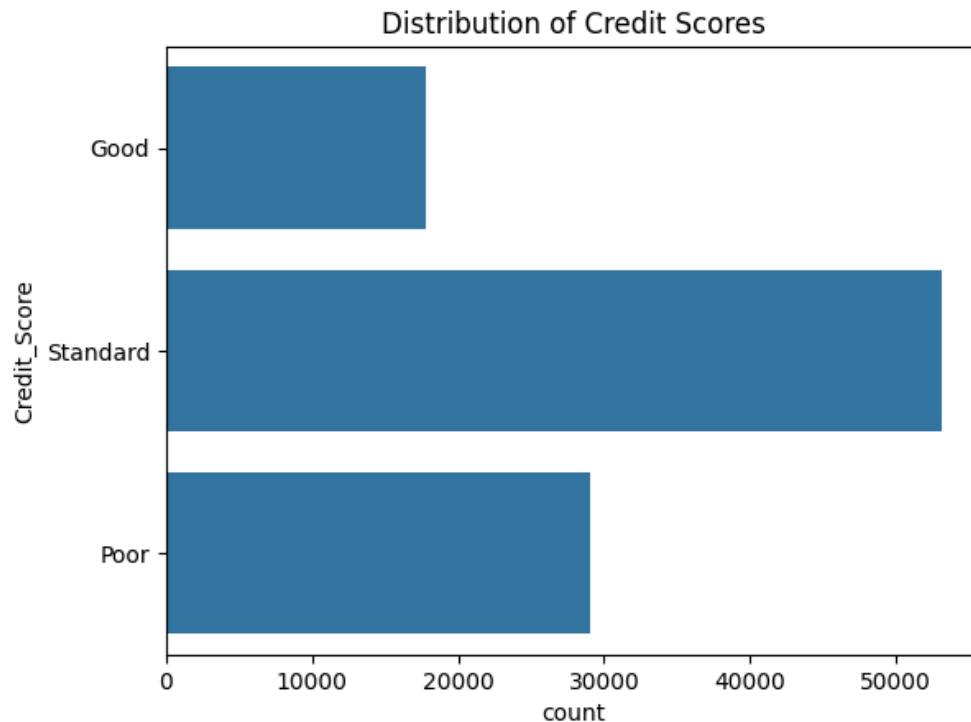
```
plt.figure(figsize=(10, 8))
```

```
sns.heatmap(numeric_data.corr(), annot=True, cmap='coolwarm')
```

```
plt.show()
```



```
# Assuming 'Credit_Score' is the target column (adjust as needed)
sns.countplot(data['Credit_Score'])
plt.title('Distribution of Credit Scores')
plt.show()
```



```
# Identify categorical columns
categorical_columns = data.select_dtypes(include=['object']).columns
print("Categorical columns:", categorical_columns)
```



```
Categorical columns: Index(['ID', 'Customer_ID', 'Month', 'Name', 'SSN', 'Occupation',
                             'Type_of_Loan', 'Credit_Mix', 'Credit_History_Age',
                             'Payment_of_Min_Amount', 'Payment_Behaviour', 'Credit_Score'],
                             dtype='object')
```

```
from sklearn.preprocessing import LabelEncoder
```

```
# Example of encoding multiple categorical columns
label_encoder = LabelEncoder()
for col in categorical_columns:
    data[col] = label_encoder.fit_transform(data[col])
```

```
from sklearn.preprocessing import LabelEncoder
```

```
# Encoding specific categorical columns
label_encoder = LabelEncoder()
for col in ['Occupation', 'Type_of_Loan', 'Credit_Mix', 'Payment_of_Min_Amount', 'Payment_Behaviour']: # Repl
    data[col] = label_encoder.fit_transform(data[col])
```

```
# Continue with your analysis and modeling
print(data.head())
```



```
   ID  Customer_ID  Month  Name  Age  SSN  Occupation  Annual_Income  \
0  16406         12320     3    84   23  10204             12      19114.12
1  16417         12320     2    84   23  10204             12      19114.12
```

2	16428	12320	6	84	23	10204	12	19114.12
3	16441	12320	0	84	23	10204	12	19114.12
4	16452	12320	7	84	23	10204	12	19114.12

	Monthly_Inhand_Salary	Num_Bank_Accounts	...	Credit_Mix	\
0	1824.843333		3	...	1
1	1824.843333		3	...	1
2	1824.843333		3	...	1
3	1824.843333		3	...	1
4	1824.843333		3	...	1

	Outstanding_Debt	Credit_Utilization_Ratio	Credit_History_Age	\
0	809.98	26.822620		186
1	809.98	31.944960		189
2	809.98	28.609352		190
3	809.98	31.377862		191
4	809.98	24.797347		192

	Payment_of_Min_Amount	Total_EMI_per_month	Amount_invested_monthly	\
0	1	49.574949		80.415295
1	1	49.574949		118.280222
2	1	49.574949		81.699521
3	1	49.574949		199.458074
4	1	49.574949		41.420153

	Payment_Behaviour	Monthly_Balance	Credit_Score
0	2	312.494089	0
1	3	284.629162	0
2	4	331.209863	0
3	5	223.451310	0
4	1	341.489231	0

[5 rows x 28 columns]

```
# Define the features (X) and target (y)
X = data.drop('Credit_Score', axis=1) # All columns except the target
y = data['Credit_Score'] # The target column
```

```
# Check the shape of X and y to confirm the split
print(f'Shape of X: {X.shape}')
print(f'Shape of y: {y.shape}')
```

```
➡ Shape of X: (100000, 27)
Shape of y: (100000,)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
# Initialize RandomForest Classifier
rf = RandomForestClassifier(random_state=42)
```

```
# Define hyperparameters grid for tuning
param_grid = {
    'n_estimators': [100, 300],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth': [4, 6, 10],
    'criterion': ['gini', 'entropy']
}
```

```
# Setup GridSearchCV for hyperparameter tuning
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, n_jobs=-1, scoring='accuracy')
```

```
# Fit the model
grid_search.fit(X_train, y_train)
```

```
# Print best parameters and best score
print(f"Best parameters: {grid_search.best_params_}")
print(f"Best accuracy: {grid_search.best_score_}")
```

➡ /usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:425: FitFailedWarning:
60 fits failed out of a total of 180.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

```
-----
37 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 729, in _fi
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 1145, in wrapper
    estimator._validate_params()
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 96, in validate
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassi
-----
23 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 729, in _fi
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 1145, in wrapper
    estimator._validate_params()
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 96, in validate
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassi

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:979: UserWarning: One or more
    nan      nan 0.68964286 0.69231429 0.68578571 0.68507143
    nan      nan 0.7306      0.73062857 0.72905714 0.72962857
    nan      nan 0.65748571 0.65747143 0.65221429 0.6519
    nan      nan 0.68828571 0.68762857 0.67322857 0.67075714
    nan      nan 0.725      0.72568571 0.72592857 0.72552857]
warnings.warn(
Best parameters: {'criterion': 'gini', 'max_depth': 10, 'max_features': 'sqrt', 'n_estimators': 300}
Best accuracy: 0.7306285714285713
```

```
# Initialize Gradient Boosting Classifier
gbc = GradientBoostingClassifier(random_state=42)

# Define hyperparameters grid for tuning
gbc_param_grid = {
    'n_estimators': [100, 150],
    'learning_rate': [0.05, 0.1],
    'max_depth': [3, 4]
}

# Setup GridSearchCV for Gradient Boosting
# gbc_grid_search = GridSearchCV(estimator=gbc, param_grid=gbc_param_grid, cv=5, n_jobs=-1, scoring='accuracy')
gbc_grid_search = GridSearchCV(
    estimator=gbc,
    param_grid=gbc_param_grid,
    cv=3, # Reduce number of folds
    n_jobs=-1,
    scoring='accuracy'
)

# Fit the model
gbc_grid_search.fit(X_train, y_train)
```

```
# Print best parameters and best score for Gradient Boosting
print(f"Best parameters for GBC: {gbc_grid_search.best_params_}")
print(f"Best accuracy for GBC: {gbc_grid_search.best_score_}")
```

```
➦ Best parameters for GBC: {'learning_rate': 0.1, 'max_depth': 4, 'n_estimators': 150}
Best accuracy for GBC: 0.729885666625124
```

```
# Predictions using the best RandomForest model
rf_best = grid_search.best_estimator_
y_pred_rf = rf_best.predict(X_test)
```

```
# Performance metrics for RandomForest
print("RandomForest Model Accuracy:", accuracy_score(y_test, y_pred_rf))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))
print("Classification Report:\n", classification_report(y_test, y_pred_rf))
```

```
➦ RandomForest Model Accuracy: 0.7298
Confusion Matrix:
[[ 4044   82 1196]
 [  678 6030 2097]
 [ 2229 1824 11820]]
Classification Report:
              precision    recall  f1-score   support

     0       0.58         0.76         0.66         5322
     1       0.76         0.68         0.72         8805
     2       0.78         0.74         0.76         15873

 accuracy          0.73         0.73         0.73         30000
 macro avg         0.71         0.73         0.71         30000
 weighted avg      0.74         0.73         0.73         30000
```

```
# Predictions using the best GradientBoosting model
gbc_best = gbc_grid_search.best_estimator_
y_pred_gbc = gbc_best.predict(X_test)
```

```
# Performance metrics for Gradient Boosting
print("Gradient Boosting Model Accuracy:", accuracy_score(y_test, y_pred_gbc))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_gbc))
print("Classification Report:\n", classification_report(y_test, y_pred_gbc))
```

```
➦ Gradient Boosting Model Accuracy: 0.7317333333333333
Confusion Matrix:
[[ 3763   86 1473]
 [  493 5982 2330]
 [ 1764 1902 12207]]
Classification Report:
              precision    recall  f1-score   support

     0       0.63         0.71         0.66         5322
     1       0.75         0.68         0.71         8805
     2       0.76         0.77         0.77         15873

 accuracy          0.73         0.73         0.73         30000
 macro avg         0.71         0.72         0.71         30000
 weighted avg      0.73         0.73         0.73         30000
```

```
from sklearn.metrics import roc_auc_score, roc_curve
```

```
# Calculate ROC-AUC using multi-class setting
rf_roc_auc = roc_auc_score(y_test, rf_best.predict_proba(X_test), multi_class='ovr')
gbc_roc_auc = roc_auc_score(y_test, gbc_best.predict_proba(X_test), multi_class='ovr')
```

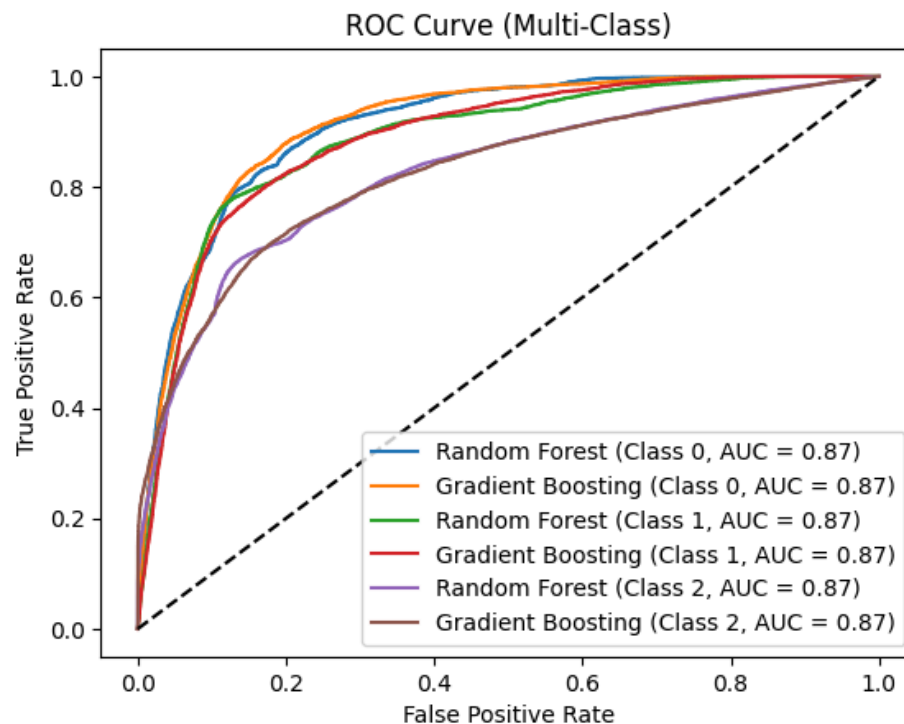
```
print(f"Random Forest ROC-AUC: {rf_roc_auc}")
print(f"Gradient Boosting ROC-AUC: {gbc_roc_auc}")
```

```
# Since you are dealing with multi-class classification, plotting ROC curves separately for each class
for i in range(len(rf_best.classes_)):
    rf_fpr, rf_tpr, _ = roc_curve(y_test, rf_best.predict_proba(X_test)[: , i], pos_label=rf_best.classes_[i])
    gbc_fpr, gbc_tpr, _ = roc_curve(y_test, gbc_best.predict_proba(X_test)[: , i], pos_label=gbc_best.classes_[i])

    plt.plot(rf_fpr, rf_tpr, label=f'Random Forest (Class {i}, AUC = {rf_roc_auc:.2f})')
    plt.plot(gbc_fpr, gbc_tpr, label=f'Gradient Boosting (Class {i}, AUC = {gbc_roc_auc:.2f})')

plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve (Multi-Class)')
plt.legend(loc='lower right')
plt.show()
```

Random Forest ROC-AUC: 0.8728681390457836
 Gradient Boosting ROC-AUC: 0.8739737432153131



Start coding or [generate](#) with AI.